

November 20, 2022

```
[ ]: import numpy as np
import pandas as pd
```

```
[ ]: df = pd.read_csv('/content/Churn_Modelling.csv')
```

```
[ ]: df
```

```
[ ]:
   RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0           1    15634602   Hargrave         619     France  Female  42
1           2    15647311     Hill         608      Spain  Female  41
2           3    15619304     Onio         502     France  Female  42
3           4    15701354     Boni         699     France  Female  39
4           5    15737888  Mitchell         850      Spain  Female  43
...         ...         ...         ...         ...         ...
9995        9996    15606229  Obijiaku         771     France    Male  39
9996        9997    15569892  Johnstone         516     France    Male  35
9997        9998    15584532      Liu         709     France  Female  36
9998        9999    15682355  Sabbatini         772  Germany    Male  42
9999       10000    15628319    Walker         792     France  Female  28
```

```

   Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0         2      0.00              1          1              1
1         1  83807.86              1          0              1
2         8 159660.80              3          1              0
3         1      0.00              2          0              0
4         2 125510.82              1          1              1
...         ...         ...         ...         ...
9995        5      0.00              2          1              0
9996       10  57369.61              1          1              1
9997        7      0.00              1          0              1
9998        3   75075.31              2          1              0
9999        4 130142.79              1          1              0
```

```

   EstimatedSalary  Exited
0         101348.88        1
1         112542.58        0
2         113931.57        1
```

3	93826.63	0
4	79084.10	0
...	...	...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

### 3. Visualization

```
[ ]: import matplotlib.pyplot as plt
```

```
[ ]: import seaborn as sns
```

```
[ ]: %matplotlib inline
```

#### (i) Univariate Analysis


```
[ ]: df[['CustomerId', 'Surname', 'CreditScore', 'Geography', 'Age', 'Tenure']].describe()
```

```
[ ]:
```

	CustomerId	CreditScore	Age	Tenure
count	1.000000e+04	10000.000000	10000.000000	10000.000000
mean	1.569094e+07	650.528800	38.921800	5.012800
std	7.193619e+04	96.653299	10.487806	2.892174
min	1.556570e+07	350.000000	18.000000	0.000000
25%	1.562853e+07	584.000000	32.000000	3.000000
50%	1.569074e+07	652.000000	37.000000	5.000000
75%	1.575323e+07	718.000000	44.000000	7.000000
max	1.581569e+07	850.000000	92.000000	10.000000

```
[ ]: sns.histplot(df.Age, kde=True)
```


```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffbaff48d10>
```



```
[ ]: # plot count plot for the gender column
sns.countplot(df.Gender)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
FutureWarning
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffbafcd90>
```



## (ii) Bivariate Analysis

```
[ ]: df[['CustomerId', 'Surname', 'CreditScore', 'Geography', 'Gender', 'Age']].corr()
```


```
[ ]:
      CustomerId  CreditScore    Age
CustomerId    1.000000    0.005308  0.009497
CreditScore    0.005308    1.000000 -0.003965
Age            0.009497   -0.003965  1.000000
```

```
[ ]: sns.scatterplot(df.CreditScore, df.Age)
      plt.ylim(0, 100)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning


```
[ ]: (0.0, 100.0)
```



(iii) Multivariate Analysis

```
[ ]: sns.pairplot(data_
↳df[['CustomerId','Geography','Gender','CreditScore','Age','Balance']],hue =_
↳'Balance')
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x7ffbaf7d2910>
```



#### 4. Descriptive Statistics

```
[ ]: #mode
df['Age'].mode()
```

```
[ ]: 0    37
      dtype: int64
```

```
[ ]: #calculation of the mean (for Age)
df["Age"].mean()
```

```
[ ]: 38.9218
```

```
[ ]: #calculation of the mean and round the result(for Age)
round(df["Age"].mean(), 2)
```

```
[ ]: 38.92
```

```
[ ]: #calculation of the median(for Age)
df["Age"].median()
```

```
[ ]: 37.0
```

```
[ ]: df.columns
```

```
[ ]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
          'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
          'IsActiveMember', 'EstimatedSalary', 'Exited'],
          dtype='object')
```

```
[ ]: df["NumOfProducts"].value_counts()
```

```
[ ]: 1    5084
     2    4590
     3     266
     4      60
     Name: NumOfProducts, dtype: int64
```

```
[ ]: df.dtypes
```

```
[ ]: RowNumber      int64
     CustomerId    int64
     Surname       object
     CreditScore   int64
     Geography     object
     Gender        object
     Age           int64
     Tenure        int64
     Balance       float64
     NumOfProducts int64
     HasCrCard     int64
     IsActiveMember int64
     EstimatedSalary float64
     Exited        int64
     dtype: object
```

```
[ ]: df.head()
```

```
[ ]:   RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0         1    15634602  Hargrave         619    France  Female   42
1         2    15647311    Hill         608    Spain  Female   41
2         3    15619304    Onio         502    France  Female   42
3         4    15701354    Boni         699    France  Female   39
4         5    15737888  Mitchell         850    Spain  Female   43
```

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
[ ]: df.describe()
```

```
[ ]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	\
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
mean	76485.889288	1.530200	0.70550	0.515100	
std	62397.405202	0.581654	0.45584	0.499797	
min	0.000000	1.000000	0.00000	0.000000	
25%	0.000000	1.000000	0.00000	0.000000	
50%	97198.540000	1.000000	1.00000	1.000000	
75%	127644.240000	2.000000	1.00000	1.000000	
max	250898.090000	4.000000	1.00000	1.000000	

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

## 5. Handling missing values



```
[ ]: df.isna().any()
```

```
[ ]: RowNumber      False
      CustomerId    False
      Surname       False
      CreditScore   False
      Geography     False
      Gender        False
      Age           False
      Tenure        False
      Balance       False
      NumOfProducts False
      HasCrCard     False
      IsActiveMember False
      EstimatedSalary False
      Exited        False
      dtype: bool
```

```
[ ]: df.isnull().sum()
```

```
[ ]: RowNumber      0
      CustomerId    0
      Surname       0
      CreditScore   0
      Geography     0
      Gender        0
      Age           0
      Tenure        0
      Balance       0
      NumOfProducts 0
      HasCrCard     0
      IsActiveMember 0
      EstimatedSalary 0
      Exited        0
      dtype: int64
```

```
[ ]: df.isnull()
```

```
[ ]:   RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0      False      False    False      False      False    False  False
1      False      False    False      False      False    False  False
2      False      False    False      False      False    False  False
3      False      False    False      False      False    False  False
4      False      False    False      False      False    False  False
...      ...      ...      ...      ...      ...      ...  ...
9995    False      False    False      False      False    False  False
9996    False      False    False      False      False    False  False
```

9997	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	False	False	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	
...	...	...	...	...	...	
9995	False	False	False	False	False	
9996	False	False	False	False	False	
9997	False	False	False	False	False	
9998	False	False	False	False	False	
9999	False	False	False	False	False	

	EstimatedSalary	Exited
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...	...	...
9995	False	False
9996	False	False
9997	False	False
9998	False	False
9999	False	False

[10000 rows x 14 columns]

```
[ ]: df.notnull()
```

```
[ ]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	True	True	True	True	True	True	True	
1	True	True	True	True	True	True	True	
2	True	True	True	True	True	True	True	
3	True	True	True	True	True	True	True	
4	True	True	True	True	True	True	True	
...	...	...	...	...	...	...	...	
9995	True	True	True	True	True	True	True	
9996	True	True	True	True	True	True	True	
9997	True	True	True	True	True	True	True	
9998	True	True	True	True	True	True	True	
9999	True	True	True	True	True	True	True	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	True	True	True	True	True	
1	True	True	True	True	True	
2	True	True	True	True	True	
3	True	True	True	True	True	
4	True	True	True	True	True	
...	...	...	...	...	...	
9995	True	True	True	True	True	
9996	True	True	True	True	True	
9997	True	True	True	True	True	
9998	True	True	True	True	True	
9999	True	True	True	True	True	


	EstimatedSalary	Exited
0	True	True
1	True	True
2	True	True
3	True	True
4	True	True
...	...	...
9995	True	True
9996	True	True
9997	True	True
9998	True	True
9999	True	True

[10000 rows x 14 columns]

## 6. Finding and replacing outliers


```
[ ]: import seaborn as sns
sns.boxplot(x=df['Age'])
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffbade3d490>
```



```
[ ]: sns.boxplot(x=df['Tenure'])
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffbaae4ac50>
```



7. Check for categorical column and perform encoding

```
[ ]: import pandas as pd
df = pd.read_csv("Churn_Modelling.csv", header=None)
```

```
[ ]: cols = df.columns
num_cols = df._get_numeric_data().columns
```

```
[ ]: num_cols
```

```
[ ]: Int64Index([], dtype='int64')
```

```
[ ]: list(set(cols) - set(num_cols))
```

```
[ ]: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13]
```

8. Split the data into dependent and independent variables

```
[ ]: # x -Independent
# y -Dependent
x = df.drop('Exited',axis=1)
y = df['Exited']
```

```
[ ]: x.head()
```

```

      □
↳ -----

      NameError                                Traceback (most recent call↳
↳ last)

      <ipython-input-53-830ed5e65d76> in <module>
      ----> 1 x.head()

      NameError: name 'x' is not defined
```

```
[ ]: y.head()
```

```

      □
↳ -----

      NameError                                Traceback (most recent call↳
↳ last)
```

```
<ipython-input-54-17b2b1f6e15b> in <module>
----> 1 y.head()
```

```
NameError: name 'y' is not defined
```

9. Scale the independent variables

```
[ ]: from sklearn import linear_model
      from sklearn.preprocessing import StandardScaler
      scale = StandardScaler()
```

```
[ ]: X = df[['Balance', 'Tenure']]
      scaledX = scale.fit_transform(X)
      print(scaledX)
```